```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")

df.head()
```

```
   customerID  gender  SeniorCitizen Partner Dependents  tenure
PhoneService  \
0  7590-VHVEG  Female              0     Yes         No       1
No
1  5575-GNVDE    Male              0      No         No      34
Yes
2  3668-QPYBK    Male              0      No         No       2
Yes
3  7795-CFOCW    Male              0      No         No      45
No
4  9237-HQITU  Female              0      No         No       2
Yes

      MultipleLines InternetService OnlineSecurity  ...
DeviceProtection  \
0  No phone service             DSL             No  ...
No
1                No             DSL            Yes  ...
Yes
2                No             DSL            Yes  ...
No
3  No phone service             DSL            Yes  ...
Yes
4                No     Fiber optic             No  ...
No

  TechSupport StreamingTV StreamingMovies        Contract
PaperlessBilling  \
0          No          No              No  Month-to-month
Yes
1          No          No              No        One year
No
2          No          No              No  Month-to-month
Yes
3         Yes          No              No        One year
No
4          No          No              No  Month-to-month
Yes

               PaymentMethod MonthlyCharges  TotalCharges Churn
0           Electronic check          29.85         29.85    No
```

```
1               Mailed check          56.95         1889.5    No
2               Mailed check          53.85          108.15   Yes
3  Bank transfer (automatic)          42.30        1840.75    No
4           Electronic check          70.70          151.65   Yes

[5 rows x 21 columns]
```

```
print("Number of rows and columns:", df.shape)
```

```
Number of rows and columns: (7043, 21)
```

```
print("\nColumn names:")
print(df.columns.tolist())
```

```
Column names:
['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn']
```

```
print("\nDataset information:")
df.info()
```

```
Dataset information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype
---  ------            --------------  -----
 0   customerID        7043 non-null   object
 1   gender            7043 non-null   object
 2   SeniorCitizen     7043 non-null   int64
 3   Partner           7043 non-null   object
 4   Dependents        7043 non-null   object
 5   tenure            7043 non-null   int64
 6   PhoneService      7043 non-null   object
 7   MultipleLines     7043 non-null   object
 8   InternetService   7043 non-null   object
 9   OnlineSecurity    7043 non-null   object
 10  OnlineBackup      7043 non-null   object
 11  DeviceProtection  7043 non-null   object
 12  TechSupport       7043 non-null   object
 13  StreamingTV       7043 non-null   object
 14  StreamingMovies   7043 non-null   object
 15  Contract          7043 non-null   object
 16  PaperlessBilling  7043 non-null   object
 17  PaymentMethod     7043 non-null   object
 18  MonthlyCharges    7043 non-null   float64
```

```
 19  TotalCharges      7043 non-null   object
 20  Churn             7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
print(df.isnull().sum())
print("\nTotal missing values:", df.isnull().sum().sum())
```

```
customerID          0
gender              0
SeniorCitizen       0
Partner             0
Dependents          0
tenure              0
PhoneService        0
MultipleLines       0
InternetService     0
OnlineSecurity      0
OnlineBackup        0
DeviceProtection    0
TechSupport         0
StreamingTV         0
StreamingMovies     0
Contract            0
PaperlessBilling    0
PaymentMethod       0
MonthlyCharges      0
TotalCharges        0
Churn               0
dtype: int64

Total missing values: 0
```

```
print("Number of duplicate rows:", df.duplicated().sum())
```

```
Number of duplicate rows: 0
```

```
print(df.dtypes)
```

```
customerID          object
gender              object
SeniorCitizen        int64
Partner             object
Dependents          object
tenure               int64
PhoneService        object
MultipleLines       object
InternetService     object
OnlineSecurity      object
OnlineBackup        object
DeviceProtection    object
```

```
TechSupport           object
StreamingTV           object
StreamingMovies       object
Contract              object
PaperlessBilling      object
PaymentMethod         object
MonthlyCharges       float64
TotalCharges          object
Churn                 object
dtype: object
```

```
print("Blank strings in TotalCharges:",
      (df["TotalCharges"].str.strip() == "").sum())

print("\nFirst 10 TotalCharges values:")
print(df["TotalCharges"].head(10).tolist())
```

```
Blank strings in TotalCharges: 11

First 10 TotalCharges values:
['29.85', '1889.5', '108.15', '1840.75', '151.65', '820.5', '1949.4',
'301.9', '3046.05', '3487.95']
```

```
df["TotalCharges"] = pd.to_numeric(
    df["TotalCharges"],
    errors="coerce"
)

print("Data type:", df["TotalCharges"].dtype)
print("Missing values in TotalCharges:",
df["TotalCharges"].isnull().sum())
```

```
Data type: float64
Missing values in TotalCharges: 11
```

```
df[df["TotalCharges"].isnull()][
    ["customerID", "tenure", "MonthlyCharges", "TotalCharges",
"Churn"]
]
```

	customerID	tenure	MonthlyCharges	TotalCharges	Churn
488	4472-LVYGI	0	52.55	NaN	No
753	3115-CZMZD	0	20.25	NaN	No
936	5709-LVOEQ	0	80.85	NaN	No
1082	4367-NUYAO	0	25.75	NaN	No
1340	1371-DWPAZ	0	56.05	NaN	No
3331	7644-OMVMY	0	19.85	NaN	No
3826	3213-VVOLG	0	25.35	NaN	No
4380	2520-SGTTA	0	20.00	NaN	No
5218	2923-ARZLG	0	19.70	NaN	No

```
6670  4075-WKNIU       0          73.35           NaN    No
6754  2775-SEFEE       0          61.90           NaN    No
```

```
df["TotalCharges"] = df["TotalCharges"].fillna(0)

print("Missing TotalCharges values:",
df["TotalCharges"].isnull().sum())
print("TotalCharges data type:", df["TotalCharges"].dtype)
```

```
Missing TotalCharges values: 0
TotalCharges data type: float64
```

```
df = df.drop("customerID", axis=1)

print("Dataset shape:", df.shape)
print("\nRemaining columns:")
print(df.columns.tolist())
```

```
Dataset shape: (7043, 20)

Remaining columns:
['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
'MonthlyCharges', 'TotalCharges', 'Churn']
```

```
print("Churn counts:")
print(df["Churn"].value_counts())

print("\nChurn percentages:")
print((df["Churn"].value_counts(normalize=True) * 100).round(2))
```

```
Churn counts:
Churn
No     5174
Yes    1869
Name: count, dtype: int64

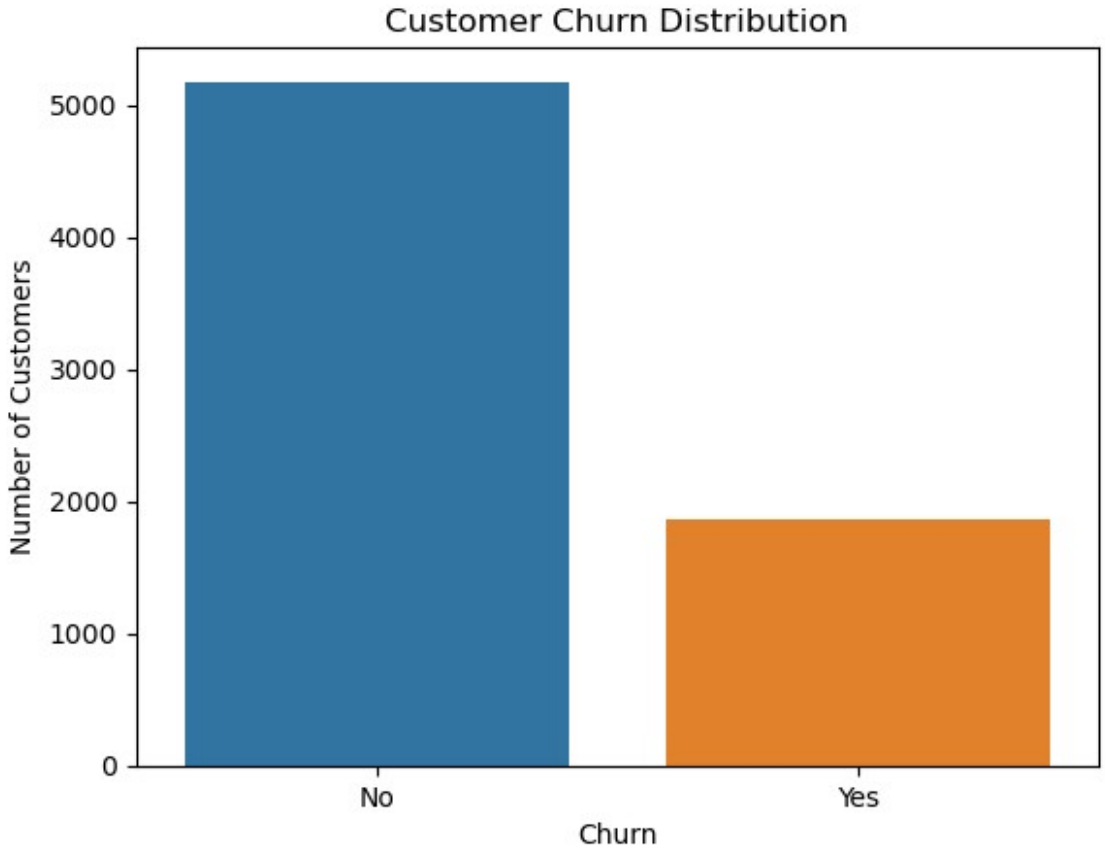
Churn percentages:
Churn
No     73.46
Yes    26.54
Name: proportion, dtype: float64
```

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.countplot(data=df, x="Churn", hue="Churn", legend=False)

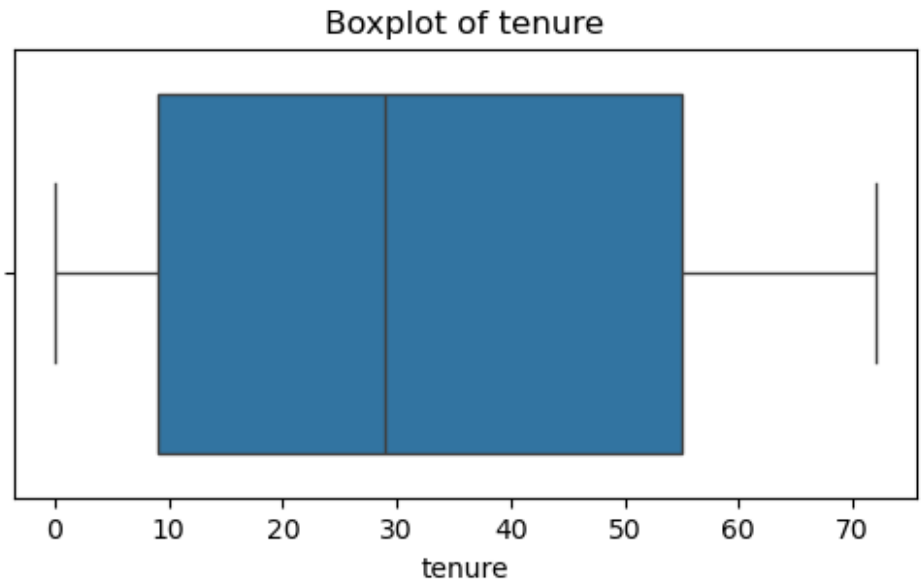
plt.title("Customer Churn Distribution")
```

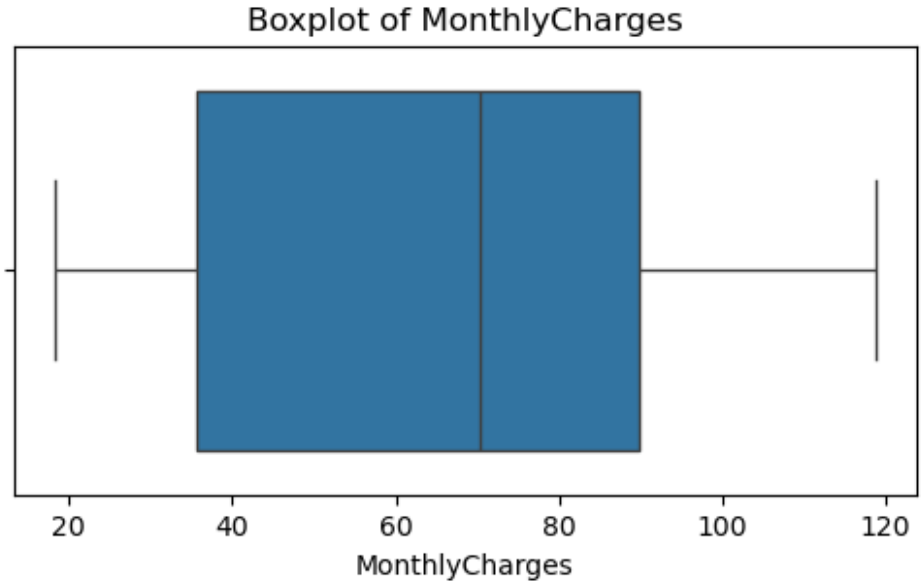
```
plt.xlabel("Churn")
plt.ylabel("Number of Customers")
plt.show()
```

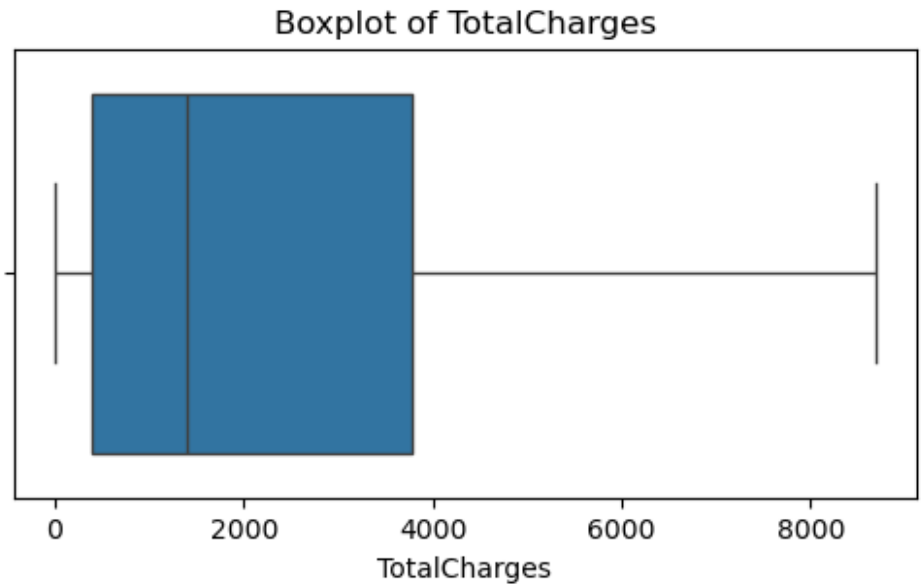



```
numerical_columns = ["tenure", "MonthlyCharges", "TotalCharges"]

for column in numerical_columns:
    plt.figure(figsize=(6, 3))
    sns.boxplot(x=df[column])
    plt.title(f"Boxplot of {column}")
    plt.xlabel(column)
    plt.show()
```


Boxplot of tenure
0
10
20
30
40
50
60
70
tenure


Boxplot of MonthlyCharges
20
40
60
80
100
120
MonthlyCharges




```
categorical_columns =
df.select_dtypes(include=["object"]).columns.tolist()
numerical_columns = df.select_dtypes(include=["int64",
"float64"]).columns.tolist()

print("Categorical columns:")
print(categorical_columns)

print("\nNumerical columns:")
print(numerical_columns)
```

```
Categorical columns:
['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
'InternetService', 'OnlineSecurity', 'OnlineBackup',
'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
'Contract', 'PaperlessBilling', 'PaymentMethod', 'Churn']

Numerical columns:
['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges']
```

```
X = df.drop("Churn", axis=1)
y = df["Churn"]

print("Features shape:", X.shape)
print("Target shape:", y.shape)

print("\nTarget values:")
print(y.value_counts())
```

```
Features shape: (7043, 19)
Target shape: (7043,)
```

```
Target values:
Churn
No     5174
Yes    1869
Name: count, dtype: int64
```

```
X_encoded = pd.get_dummies(X, drop_first=True, dtype=int)

# Encode the target variable
y_encoded = y.map({"No": 0, "Yes": 1})

print("Original feature shape:", X.shape)
print("Encoded feature shape:", X_encoded.shape)

print("\nFirst 5 rows of encoded features:")
print(X_encoded.head())

print("\nEncoded target values:")
print(y_encoded.value_counts())
```

```
Original feature shape: (7043, 19)
Encoded feature shape: (7043, 30)

First 5 rows of encoded features:
   SeniorCitizen  tenure  MonthlyCharges  TotalCharges  gender_Male  \
0              0       1           29.85         29.85            0
1              0      34           56.95       1889.50            1
2              0       2           53.85        108.15            1
3              0      45           42.30       1840.75            1
4              0       2           70.70        151.65            0

   Partner_Yes  Dependents_Yes  PhoneService_Yes  \
0            1               0                 0
1            0               0                 1
2            0               0                 1
3            0               0                 0
4            0               0                 1

   MultipleLines_No phone service  MultipleLines_Yes  ...  \
0                               1                  0  ...
1                               0                  0  ...
2                               0                  0  ...
3                               1                  0  ...
4                               0                  0  ...

   StreamingTV_No internet service  StreamingTV_Yes  \
0                                0                0
1                                0                0
2                                0                0
3                                0                0
```

```
4                                     0                  0

   StreamingMovies_No internet service  StreamingMovies_Yes  \
0                                    0                    0
1                                    0                    0
2                                    0                    0
3                                    0                    0
4                                    0                    0

   Contract_One year  Contract_Two year  PaperlessBilling_Yes  \
0                  0                  0                     1
1                  1                  0                     0
2                  0                  0                     1
3                  1                  0                     0
4                  0                  0                     1

   PaymentMethod_Credit card (automatic)  PaymentMethod_Electronic
check  \
0                                      0
1
1                                      0
0
2                                      0
0
3                                      0
0
4                                      0
1

   PaymentMethod_Mailed check
0                           0
1                           1
2                           1
3                           0
4                           0

[5 rows x 30 columns]

Encoded target values:
Churn
0    5174
1    1869
Name: count, dtype: int64

from sklearn.model_selection import train_test_split

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded,
    y_encoded,
```

```
    test_size=0.2,
    random_state=42,
    stratify=y_encoded
)

print("Training features shape:", X_train.shape)
print("Testing features shape:", X_test.shape)

print("\nTraining target shape:", y_train.shape)
print("Testing target shape:", y_test.shape)

print("\nChurn distribution in training data:")
print(y_train.value_counts(normalize=True).mul(100).round(2))

print("\nChurn distribution in testing data:")
print(y_test.value_counts(normalize=True).mul(100).round(2))
```

```
Training features shape: (5634, 30)
Testing features shape: (1409, 30)

Training target shape: (5634,)
Testing target shape: (1409,)

Churn distribution in training data:
Churn
0    73.46
1    26.54
Name: proportion, dtype: float64

Churn distribution in testing data:
Churn
0    73.46
1    26.54
Name: proportion, dtype: float64
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Scaled training data shape:", X_train_scaled.shape)
print("Scaled testing data shape:", X_test_scaled.shape)
```

```
Scaled training data shape: (5634, 30)
Scaled testing data shape: (1409, 30)
```

```
from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression(random_state=42, max_iter=1000)
lr_model.fit(X_train_scaled, y_train)
y_pred_lr = lr_model.predict(X_test_scaled)
```

```
print("Logistic Regression model trained successfully!")
print("First 10 actual values:")
print(y_test.iloc[:10].tolist())

print("\nFirst 10 predicted values:")
print(y_pred_lr[:10].tolist())
```

```
Logistic Regression model trained successfully!
First 10 actual values:
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1]

First 10 predicted values:
[0, 1, 0, 0, 0, 1, 0, 0, 0, 0]
```

```
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
    classification_report,
    ConfusionMatrixDisplay
)
accuracy_lr = accuracy_score(y_test, y_pred_lr)
precision_lr = precision_score(y_test, y_pred_lr)
recall_lr = recall_score(y_test, y_pred_lr)
f1_lr = f1_score(y_test, y_pred_lr)

print("Logistic Regression Evaluation")
print("--------------------------------")
print("Accuracy:", round(accuracy_lr, 4))
print("Precision:", round(precision_lr, 4))
print("Recall:", round(recall_lr, 4))
print("F1-score:", round(f1_lr, 4))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_lr))

cm_lr = confusion_matrix(y_test, y_pred_lr)

print("Confusion Matrix:")
print(cm_lr)

ConfusionMatrixDisplay(
    confusion_matrix=cm_lr,
    display_labels=["No Churn", "Churn"]
).plot(cmap="Blues")

plt.title("Logistic Regression Confusion Matrix")
plt.show()
```

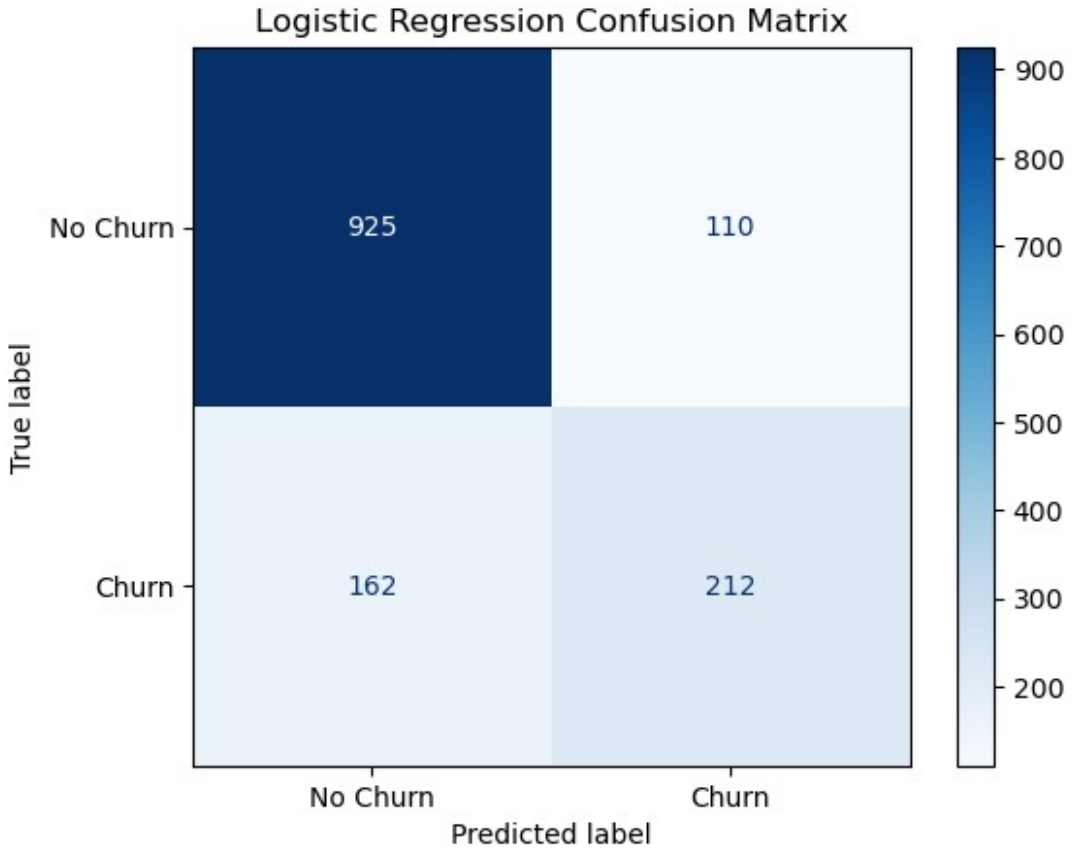
```
Logistic Regression Evaluation
--------------------------------
Accuracy: 0.807
Precision: 0.6584
Recall: 0.5668
F1-score: 0.6092

Classification Report:
              precision    recall  f1-score   support

           0       0.85      0.89      0.87      1035
           1       0.66      0.57      0.61       374

    accuracy                           0.81      1409
   macro avg       0.75      0.73      0.74      1409
weighted avg       0.80      0.81      0.80      1409

Confusion Matrix:
[[925 110]
 [162 212]]
```




```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_model = DecisionTreeClassifier(random_state=42)

dt_model.fit(X_train, y_train)

y_pred_dt = dt_model.predict(X_test)

print("Decision Tree model trained successfully!")

print("\nFirst 10 actual values:")
print(y_test.iloc[:10].tolist())

print("\nFirst 10 predicted values:")
print(y_pred_dt[:10].tolist())
```

```
Decision Tree model trained successfully!

First 10 actual values:
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1]

First 10 predicted values:
[0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
```

```
accuracy_dt = accuracy_score(y_test, y_pred_dt)
precision_dt = precision_score(y_test, y_pred_dt)
recall_dt = recall_score(y_test, y_pred_dt)
f1_dt = f1_score(y_test, y_pred_dt)

print("Decision Tree Evaluation")
print("------------------------")
print("Accuracy:", round(accuracy_dt, 4))
print("Precision:", round(precision_dt, 4))
print("Recall:", round(recall_dt, 4))
print("F1-score:", round(f1_dt, 4))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_dt))

cm_dt = confusion_matrix(y_test, y_pred_dt)

print("Confusion Matrix:")
print(cm_dt)

ConfusionMatrixDisplay(
    confusion_matrix=cm_dt,
    display_labels=["No Churn", "Churn"]
).plot(cmap="Blues")

plt.title("Decision Tree Confusion Matrix")
plt.show()
```

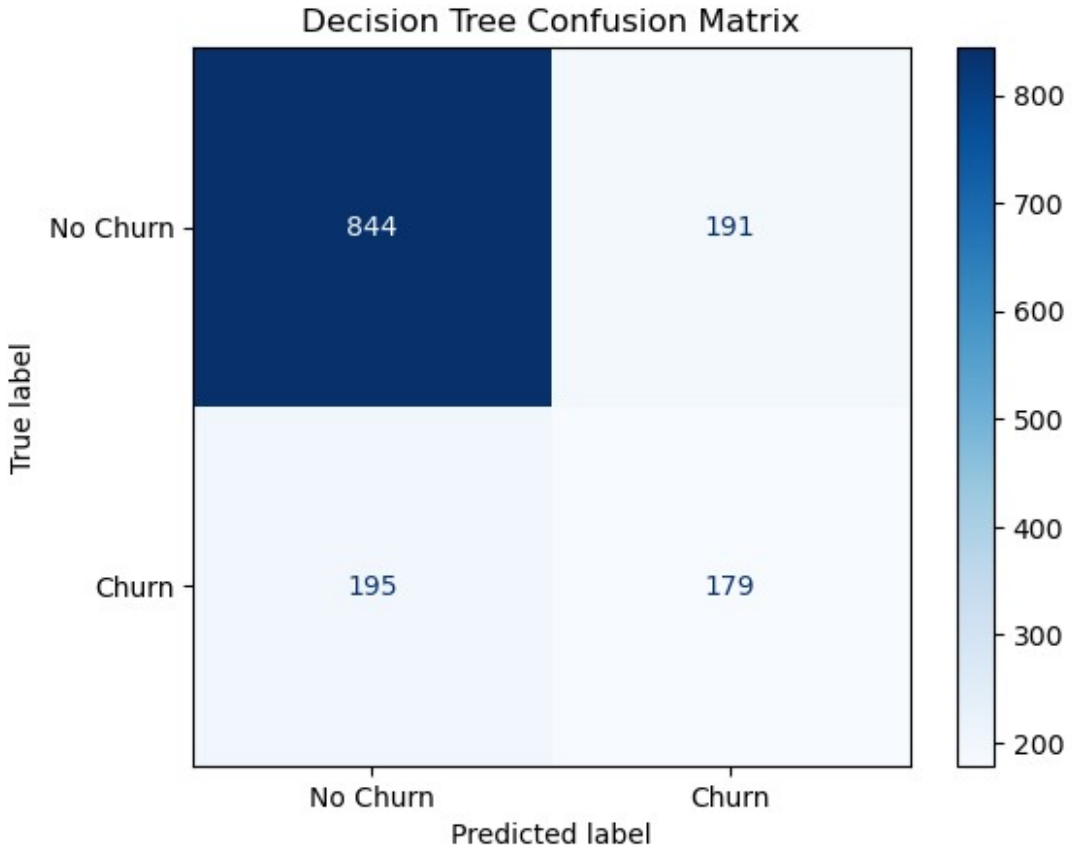
```
Decision Tree Evaluation
------------------------
Accuracy: 0.726
Precision: 0.4838
Recall: 0.4786
F1-score: 0.4812

Classification Report:
              precision    recall  f1-score   support

           0       0.81      0.82      0.81      1035
           1       0.48      0.48      0.48       374

    accuracy                           0.73      1409
   macro avg       0.65      0.65      0.65      1409
weighted avg       0.73      0.73      0.73      1409

Confusion Matrix:
[[844 191]
 [195 179]]
```




```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_model = RandomForestClassifier(
    n_estimators=100,
    random_state=42
)

rf_model.fit(X_train, y_train)

y_pred_rf = rf_model.predict(X_test)

print("Random Forest model trained successfully!")

print("\nFirst 10 actual values:")
print(y_test.iloc[:10].tolist())

print("\nFirst 10 predicted values:")
print(y_pred_rf[:10].tolist())
```

```
Random Forest model trained successfully!

First 10 actual values:
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1]

First 10 predicted values:
[0, 1, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)

print("Random Forest Evaluation")
print("------------------------")
print("Accuracy:", round(accuracy_rf, 4))
print("Precision:", round(precision_rf, 4))
print("Recall:", round(recall_rf, 4))
print("F1-score:", round(f1_rf, 4))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf))

cm_rf = confusion_matrix(y_test, y_pred_rf)

print("Confusion Matrix:")
print(cm_rf)

ConfusionMatrixDisplay(
    confusion_matrix=cm_rf,
    display_labels=["No Churn", "Churn"]
).plot(cmap="Blues")
```

```
plt.title("Random Forest Confusion Matrix")
plt.show()

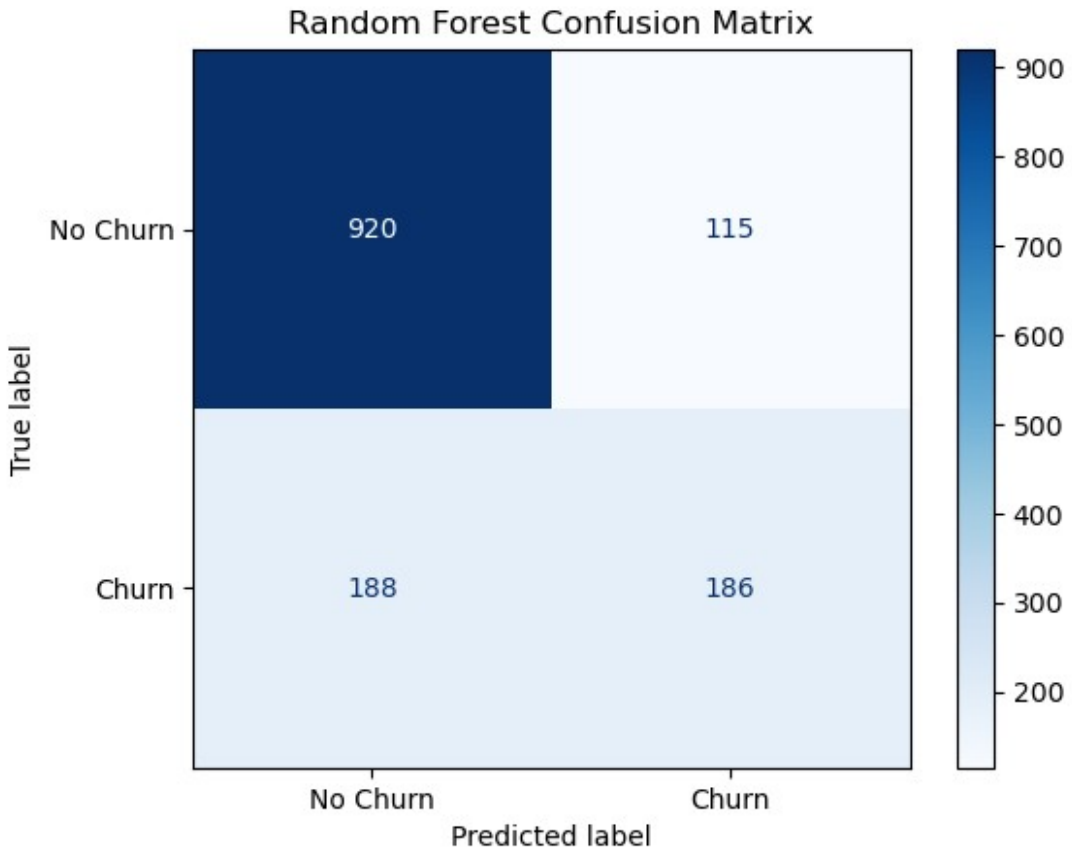
Random Forest Evaluation
------------------------
Accuracy: 0.785
Precision: 0.6179
Recall: 0.4973
F1-score: 0.5511

Classification Report:
              precision    recall  f1-score   support

           0       0.83      0.89      0.86      1035
           1       0.62      0.50      0.55       374

    accuracy                           0.78      1409
   macro avg       0.72      0.69      0.70      1409
weighted avg       0.77      0.78      0.78      1409

Confusion Matrix:
[[920 115]
 [188 186]]
```



```
results = pd.DataFrame({
    "Model": [
        "Logistic Regression",
        "Decision Tree",
        "Random Forest"
    ],
    "Accuracy": [
        accuracy_lr,
        accuracy_dt,
        accuracy_rf
    ],
    "Precision": [
        precision_lr,
        precision_dt,
        precision_rf
    ],
    "Recall": [
        recall_lr,
        recall_dt,
        recall_rf
    ],
    "F1-score": [
        f1_lr,
        f1_dt,
        f1_rf
    ]
})

print("Model Performance Comparison:")
print(results.round(4).to_string(index=False))
```

```
Model Performance Comparison:
              Model  Accuracy  Precision  Recall  F1-score
Logistic Regression     0.807     0.6584  0.5668    0.6092
      Decision Tree     0.726     0.4838  0.4786    0.4812
      Random Forest     0.785     0.6179  0.4973    0.5511
```

```
feature_importance = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": rf_model.feature_importances_
})

feature_importance = feature_importance.sort_values(
    by="Importance",
    ascending=False
)

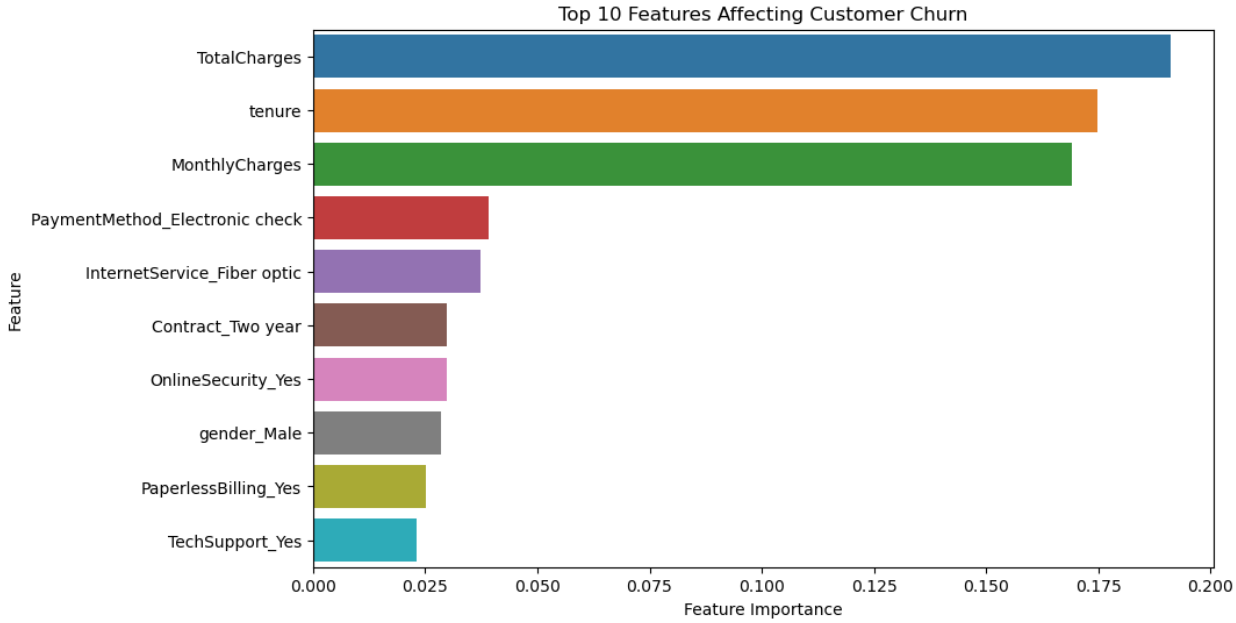
print("Top 10 Important Features:")
print(feature_importance.head(10).to_string(index=False))
```

```
top_features = feature_importance.head(10)

plt.figure(figsize=(10, 6))
sns.barplot(
    data=top_features,
    x="Importance",
    y="Feature",
    hue="Feature",
    legend=False
)

plt.title("Top 10 Features Affecting Customer Churn")
plt.xlabel("Feature Importance")
plt.ylabel("Feature")
plt.show()
```

```
Top 10 Important Features:
                          Feature  Importance
                    TotalCharges    0.191149
                          tenure    0.174795
                  MonthlyCharges    0.169026
PaymentMethod_Electronic check    0.039095
   InternetService_Fiber optic    0.037259
              Contract_Two year    0.029864
             OnlineSecurity_Yes    0.029774
                    gender_Male    0.028501
           PaperlessBilling_Yes    0.025097
                TechSupport_Yes    0.023192
```



Business Conclusion

The customer churn prediction project was completed using Logistic Regression, Decision Tree, and Random Forest classification algorithms. Among the three models, Logistic Regression performed best, achieving an accuracy of **80.70%** and an F1-score of **60.92%**. The feature importance analysis identified TotalCharges, tenure, and MonthlyCharges as the three most important predictive features.

The telecommunications company can use these findings to identify customers who may be at risk of leaving and develop targeted retention strategies, such as personalized offers, contract incentives, and improved customer support. However, further improvement in churn detection is needed because the best model identified only 56.68% of actual churners.